

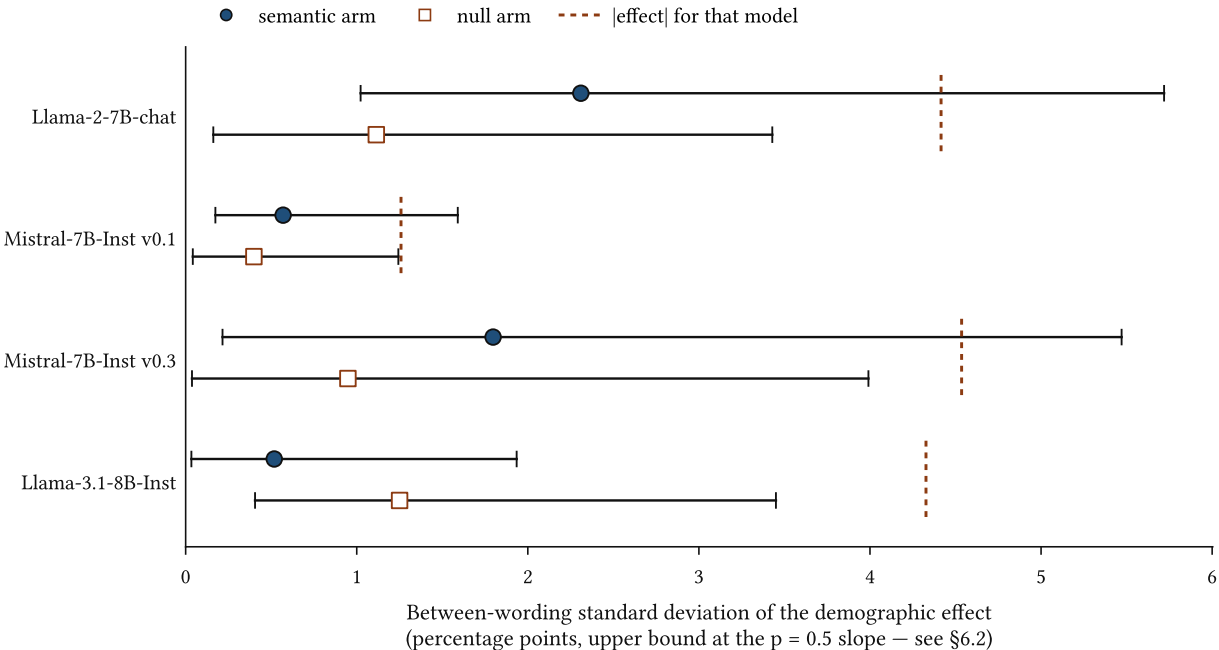


Figure 5. Wording moves the demographic effect by 0.11 to 0.53 of the effect itself in point estimate, and surface form does as much work as meaning. Posterior median and 95% credible interval for the between-wording standard deviation from a crossed random-effects model, fitted separately to the six semantic paraphrases and the six semantically null perturbations. Sampling noise is partialled out by partial pooling, so this is dispersion rather than dispersion-plus-noise. The dashed rule marks the magnitude of that model's own demographic effect. Percentage points are computed at the logistic's maximum slope and are therefore upper bounds (§6.2). The dispersion-to-effect ratio itself is poorly pinned: on Mistral-7B-Inst v0.1's semantic arm its 95% interval runs [0.08, 12.37]. This figure orders nothing; it shows that the component is large and that the two arms overlap.