

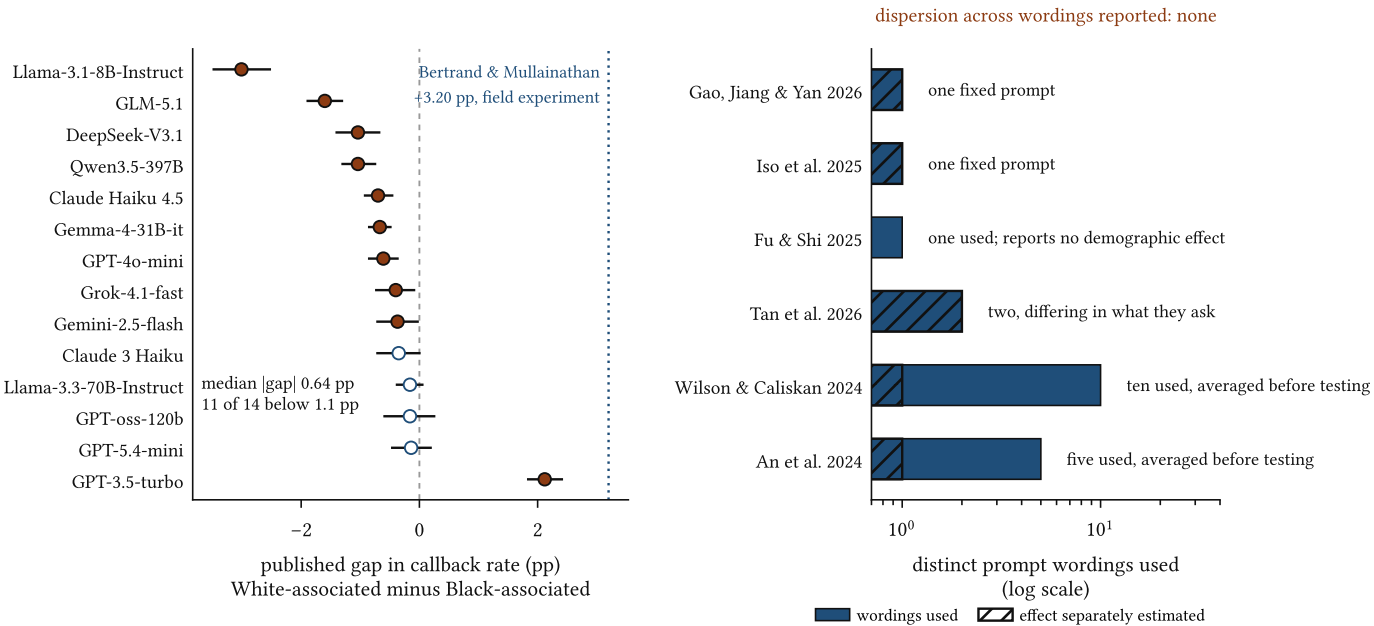


Figure 8. The field's effects are small, and none of them is published with a range. Left: every demographic gap in the largest published model panel (Gao, Jiang and Yan, Table I), with 95% intervals; filled points are significant. The dotted line is the Bertrand and Mullainathan field experiment, the anchor this literature compares itself to. The median published gap is 0.64 points. Right: how many distinct wordings each study estimated its effect under. two use one prompt, one uses two, and two average over many and publish only the mean, which discards the dispersion rather than measuring it. one further study reports no demographic effect, so no wording count applies to it and its hatched overlay is omitted. Our own effect is on a log-odds scale and is deliberately not plotted on the left axis, because converting it to points requires a base rate these models do not share. The quantity that does transfer is the ratio of wording-induced dispersion to effect size, which on the two of our models whose effect is distinguishable from zero runs from 25 to 33 per cent, and on the other two reaches 118 per cent against a denominator that is not separable from zero.