




Figure 6. The four models do not share an operating point, so they do not share a percentage-point scale. Left: the logistic curve, the tangent an audit assumes when it converts at $p = 0.5$ (grey, slope 0.25), and the tangent that is actually there at each model's mean operating point (blue). The gap between the two slopes is the conversion error. Right: that error, 0.25 divided by the model's own mean $p(1-p)$, on a log axis. It ranges over more than two orders of magnitude across four checkpoints of the same size measured on the same design. The quantity has no effect size in it, so it is defined whether or not the model has an effect to convert; the sienna tick marks the overstatement realised on a measured effect, shown only for the two models whose effect is distinguishable from zero, where that quotient is stable.